

A Summary of “A Simple Framework for Contrastive Learning of Visual Representations”

Jasmine Nelson

A Simple Framework for Contrastive Learning of Visual Representations is a paper authored by Chen, Kornblith, Norouzi, and Hinton. In this paper, the researchers propose a new framework designed for contrastive learning of visual representations. The method introduced in the paper is SimCLR and the framework simplifies existing self-supervised learning algorithms without the need for “specialized architectures or a memory bank” (Chen, Kornblith, Norouzi, & Hinton, 2020).

Contrastive learning is a technique used in machine learning. Contrastive learning involves comparing images to generate a similarity score. The method clusters similar representations, or positive pairs, together while maximizing the distance between dissimilar representations, also known as negative pairs. For example, an image of a cat and a dog would receive a higher similarity score than an image of a cat and a car. The higher similarity score indicates that the visual representations are more similar.

“Learning effective visual representations without human supervision is a long-standing problem. (Chen et al., 2020)” As stated in the paper, most commonly used techniques are either generative or discriminative. Generative techniques use pixel-level generation which is computationally steep. However, progress is being made with discriminative techniques based on contrastive learning. Although the work of learning effective visual representations has been advanced through discriminative techniques based on contrastive learning, the state-of-the-art recommendations prior to this research include using specific architectures and memory banks to support the model. However, in the new approach detailed in the paper, the researchers propose a new method that simplifies the existing recommendations and that does not require specific architectures or a memory bank.

To develop a method that simplifies the existing recommendations, the researchers wanted to understand what allows the contrastive predictive tasks to learn representations. Achieving this understanding required the researchers to study the major components of their framework. They found that the model’s representation learning improved when they combined multiple data augmentations, used a learnable nonlinear transformation between the contrastive loss and the representation, normalized embeddings and performed temperature adjustments, and included larger batch sizes and longer training. The researchers combined all of these findings to generate a new state-of-the-art framework.

Combining Multiple Data Augmentations

The SimCLR framework learns representations by comparing samples with different augmentations and trying to find the maximum agreement between the different views. The researchers applied 3 augmentations sequentially: random cropping of the samples, random color distortions, and random Gaussian blur. The researchers found that random cropping alone does not increase performance because most patches from the same image share similar colors. Therefore, better performance depends on using cropping alongside the color distortions. Furthermore, they found that stronger color augmentation resulted in a substantial improvement especially in unsupervised contrastive learning.

Using a learnable nonlinear transformation between the contrastive loss and the representation

A major component of the SimCLR framework includes the use of a multi-layer perceptron (MLP) with one hidden layer and ReLU non-linearity. This component was developed as a result of the researchers studying the effects of including a projection head. The researchers examined the results of using three different architectures: identity mapping, linear projection, and nonlinear projection with 1 additional hidden layer and ReLU activation. They found that using nonlinear projection improves the quality of the learned representations better than a linear projection and much better than using no projection at all. Furthermore, since they trained the encoder network and the projection head to maximize agreement using the contrastive loss, they were able to remove the projection head after training and only use the encoder and representation for downstream tasks.

Normalized embeddings and temperature adjustments

Normalized embeddings and temperature adjustments are critical components of the SimCLR framework. To gather this conclusion, the researchers randomly sampled minibatches of images. For each minibatch, the researchers choose an example image and augmented it with some of the prior mentioned data-augmentation techniques (random cropping, color distortions, and/or Gaussian blur). The augmented image was included in the minibatch. The researchers then used the contrastive prediction task on the sample image to find the contrastive loss between the sample image and the remaining images in the batch. The loss function rewards the model with a lower loss when there is a high

similarity score between the positive pair and penalizes the model with a higher loss when it shows similarity between any of the negative pairs. To score this loss, the researchers used the normalized temperature-scaled cross entropy (NT-Xent) loss. The researchers decided to use this loss function after comparing it to other common loss functions, such as logistic loss and margin loss. They found the combination of ℓ_2 normalization with a tunable temperature parameter is significantly more effective at helping the model learn especially in the case of hard negatives.

Larger batch sizes and longer training

The final component of the SimCLR framework involves using larger batch sizes and longer training times. The researchers found that using larger batch sizes and longer training provides more negative examples, allowing them to use fewer epochs to reach a specific accuracy. Overall, they found that the boost in performance is significant when the number of training epochs is kept fairly small, around 100.

In conclusion, although the SimCLR framework is novel in its implementation, the researchers declare that it's merely a combination of previously used methods. However, in the field of machine learning, combining techniques to create new, more effective frameworks can provide cutting edge results. "A Simple Framework for Contrastive Learning of Visual Representations" provides the field with a novel implementation that provides a new state-of-the-art way to learn effective visual representations without human supervision which certainly represents a meaningful contribution to the field.

References

Chen, T., Kornblith, S., Norouzi, M., & Hinton, G. (2020). A Simple Framework for Contrastive Learning of Visual Representations. *Proceedings of the 37th International Conference on Machine Learning*.